

```
In [2]: import pickle
import matplotlib.pyplot as plt
import pandas as pd
import numpy as np
import seaborn as sns
from sklearn.preprocessing import MinMaxScaler
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.metrics import ConfusionMatrixDisplay , classification_report
from sklearn.model_selection import train_test_split
```

```
In [4]: df_train = pd.read_csv("train.csv")
df_train.T
```

```
Out[4]:
```

	0	1	2	3	4	5	6	7	8	9	...
battery_power	842.0	1021.0	563.0	615.0	1821.0	1859.0	1821.0	1954.0	1445.0	509.0	...
blue	0.0	1.0	1.0	1.0	1.0	0.0	0.0	0.0	1.0	1.0	...
clock_speed	2.2	0.5	0.5	2.5	1.2	0.5	1.7	0.5	0.5	0.6	...
dual_sim	0.0	1.0	1.0	0.0	0.0	1.0	0.0	1.0	0.0	1.0	...
fc	1.0	0.0	2.0	0.0	13.0	3.0	4.0	0.0	0.0	2.0	...
four_g	0.0	1.0	1.0	0.0	1.0	0.0	1.0	0.0	0.0	1.0	...
int_memory	7.0	53.0	41.0	10.0	44.0	22.0	10.0	24.0	53.0	9.0	...
m_dep	0.6	0.7	0.9	0.8	0.6	0.7	0.8	0.8	0.7	0.1	...
mobile_wt	188.0	136.0	145.0	131.0	141.0	164.0	139.0	187.0	174.0	93.0	...
n_cores	2.0	3.0	5.0	6.0	2.0	1.0	8.0	4.0	7.0	5.0	...
pc	2.0	6.0	6.0	9.0	14.0	7.0	10.0	0.0	14.0	15.0	...
px_height	20.0	905.0	1263.0	1216.0	1208.0	1004.0	381.0	512.0	386.0	1137.0	...
px_width	756.0	1988.0	1716.0	1786.0	1212.0	1654.0	1018.0	1149.0	836.0	1224.0	...
ram	2549.0	2631.0	2603.0	2769.0	1411.0	1067.0	3220.0	700.0	1099.0	513.0	...
sc_h	9.0	17.0	11.0	16.0	8.0	17.0	13.0	16.0	17.0	19.0	...
sc_w	7.0	3.0	2.0	8.0	2.0	1.0	8.0	3.0	1.0	10.0	...
talk_time	19.0	7.0	9.0	11.0	15.0	10.0	18.0	5.0	20.0	12.0	...
three_g	0.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	...
touch_screen	0.0	1.0	1.0	0.0	1.0	0.0	0.0	1.0	0.0	0.0	...
wifi	1.0	0.0	0.0	0.0	0.0	0.0	1.0	1.0	0.0	0.0	...
price_range	1.0	2.0	2.0	2.0	1.0	1.0	3.0	0.0	0.0	0.0	...

21 rows × 2000 columns



In [5]: df_train.info()


```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2000 entries, 0 to 1999
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   battery_power          2000 non-null   int64
1   blue                   2000 non-null   int64
2   clock_speed            2000 non-null   float64
3   dual_sim               2000 non-null   int64
4   fc                     2000 non-null   int64
5   four_g                 2000 non-null   int64
6   int_memory             2000 non-null   int64
7   m_dep                  2000 non-null   float64
8   mobile_wt              2000 non-null   int64
9   n_cores                2000 non-null   int64
10  pc                     2000 non-null   int64
11  px_height               2000 non-null   int64
12  px_width               2000 non-null   int64
13  ram                    2000 non-null   int64
14  sc_h                   2000 non-null   int64
15  sc_w                   2000 non-null   int64
16  talk_time              2000 non-null   int64
17  three_g                2000 non-null   int64
18  touch_screen           2000 non-null   int64
19  wifi                   2000 non-null   int64
20  price_range            2000 non-null   int64
dtypes: float64(2), int64(19)
memory usage: 328.3 KB
```

In [8]: df_train.describe()

Out[8]:

	battery_power	blue	clock_speed	dual_sim	fc	four_g	int_memory
count	2000.000000	2000.0000	2000.000000	2000.000000	2000.000000	2000.000000	2000.000000
mean	1238.518500	0.4950	1.522250	0.509500	4.309500	0.521500	32.046500
std	439.418206	0.5001	0.816004	0.500035	4.341444	0.499662	18.145710
min	501.000000	0.0000	0.500000	0.000000	0.000000	0.000000	2.000000
25%	851.750000	0.0000	0.700000	0.000000	1.000000	0.000000	16.000000
50%	1226.000000	0.0000	1.500000	1.000000	3.000000	1.000000	32.000000
75%	1615.250000	1.0000	2.200000	1.000000	7.000000	1.000000	48.000000
max	1998.000000	1.0000	3.000000	1.000000	19.000000	1.000000	64.000000

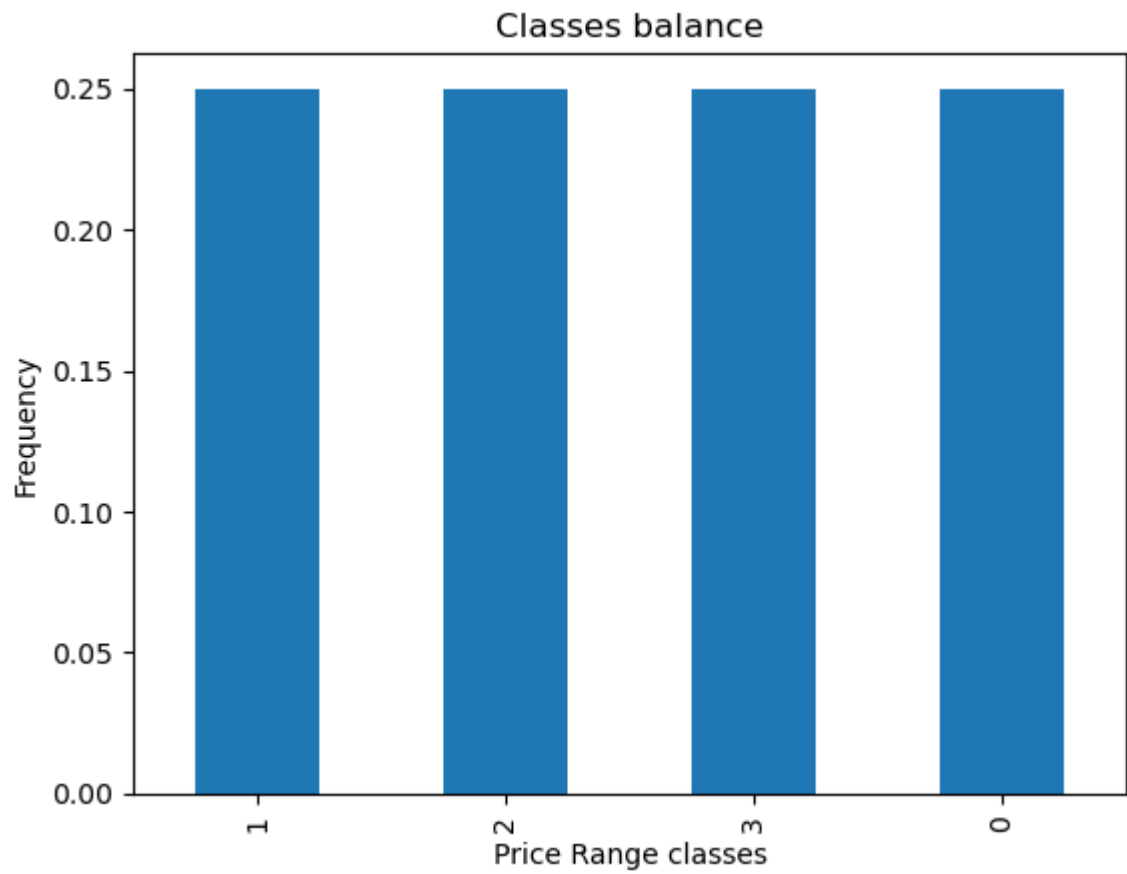
8 rows × 21 columns



```
In [9]: df_train.nunique()
```

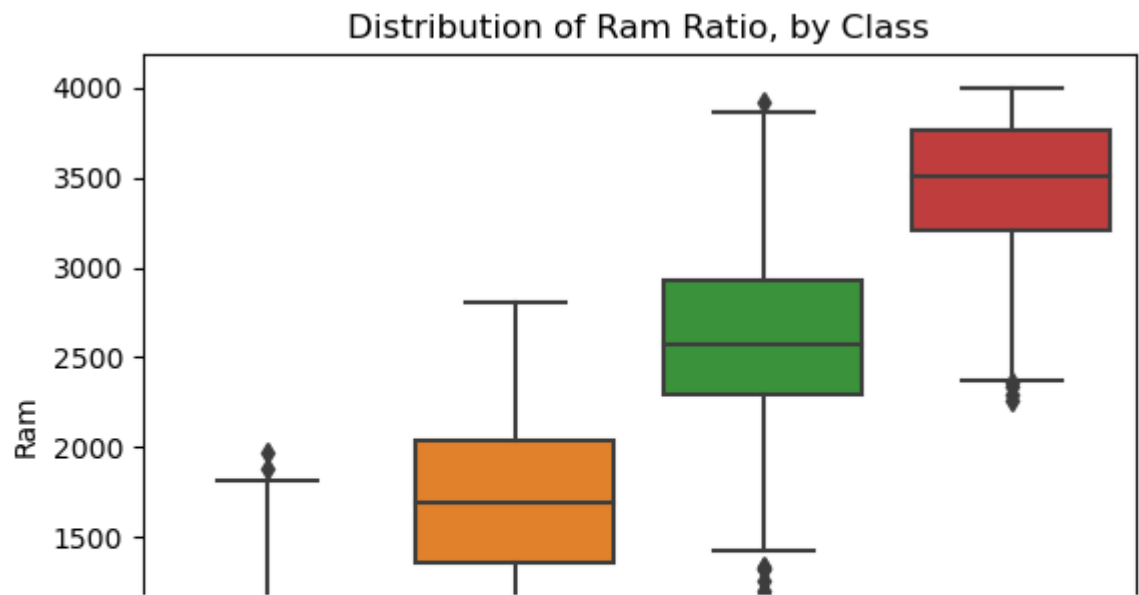
```
Out[9]: battery_power    1094  
blue                    2  
clock_speed            26  
dual_sim               2  
fc                     20  
four_g                 2  
int_memory             63  
m_dep                  10  
mobile_wt              121  
n_cores                8  
pc                     21  
px_height              1137  
px_width               1109  
ram                    1562  
sc_h                   15  
sc_w                   19  
talk_time              19  
three_g                2  
touch_screen           2  
wifi                   2  
price_range            4  
dtype: int64
```

```
In [11]: df_train['price_range'].value_counts(normalize=True).plot(kind = 'bar')
plt.xlabel("Price Range classes")
plt.ylabel("Frequency")
plt.title("Classes balance");
```

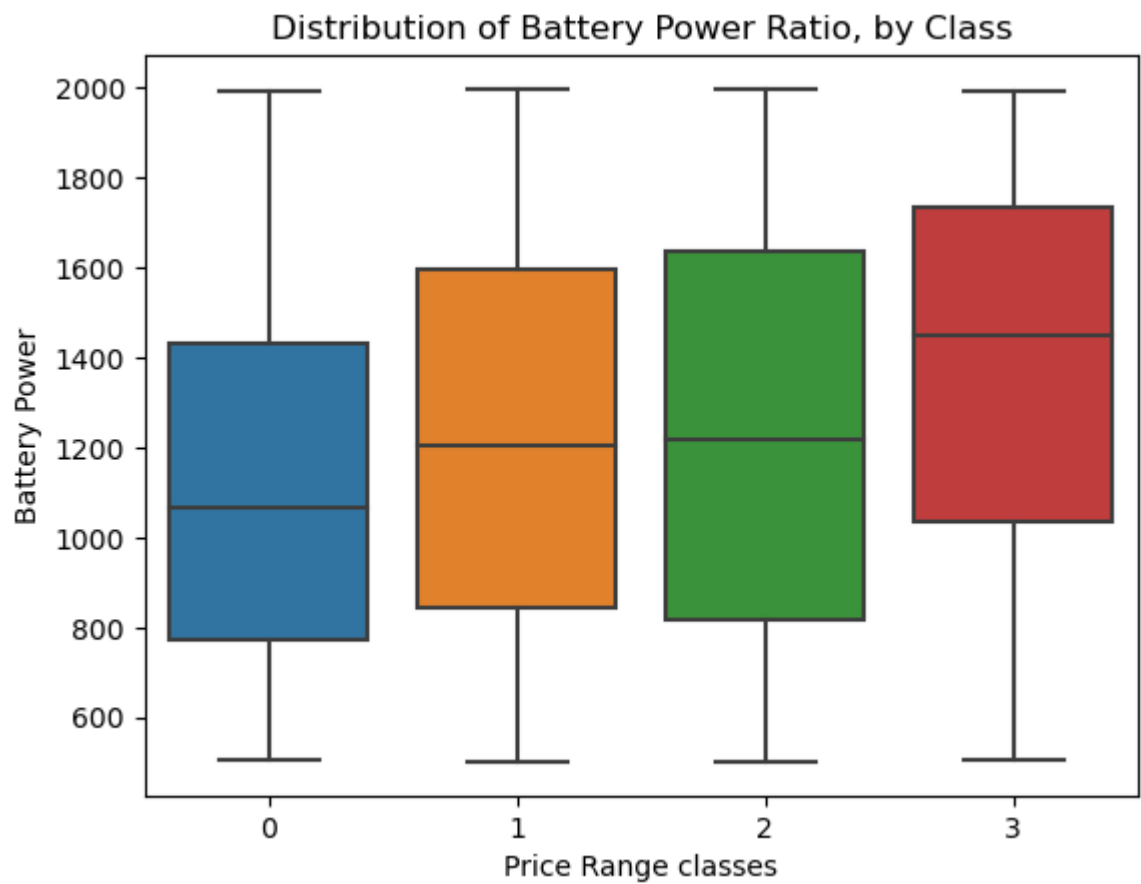


```
In [14]: sns.boxplot(x= 'price_range', y='ram', data=df_train)
plt.xlabel("Price Range classes")
plt.ylabel("Ram")
plt.title("Distribution of Ram Ratio, by Class")
```


```
Out[14]: Text(0.5, 1.0, 'Distribution of Ram Ratio, by Class')
```




```
In [16]: sns.boxplot(x= 'price_range', y='battery_power', data=df_train)
plt.xlabel("Price Range classes")
plt.ylabel("Battery Power")
plt.title("Distribution of Battery Power Ratio, by Class");
```



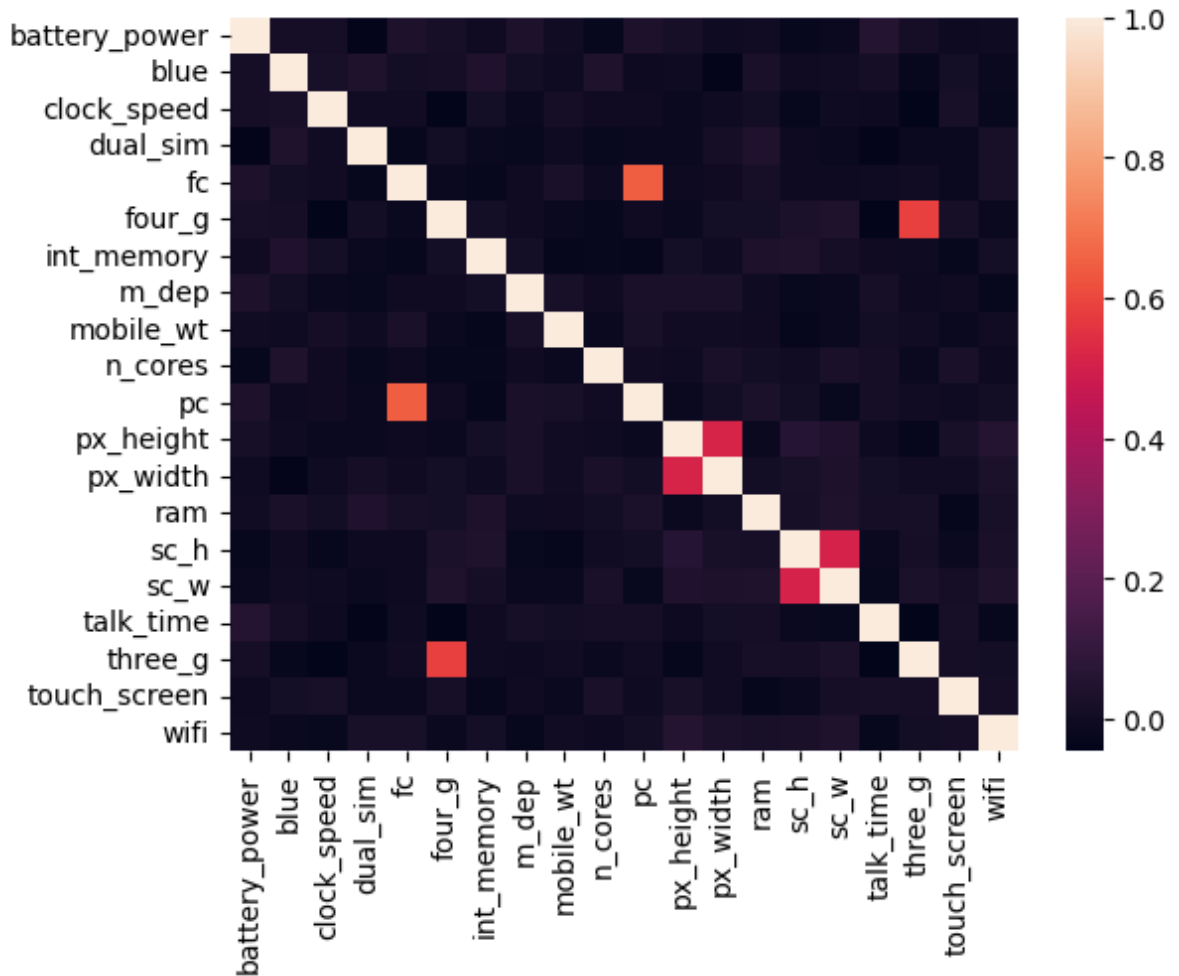
```
In [20]: labels = ["3G-supported", 'Not supported']  
values = df_train['three_g'].value_counts().values  
fig, ax = plt.subplots()  
colors = ['orange', 'lightskyblue']  
ax.pie(values, labels=labels, autopct='%1.1f%%', shadow=True, startangle=90, color  
plt.show();
```



```
In [21]: labels = ["4G-supported", 'Not supported']  
values = df_train['four_g'].value_counts().values  
fig1, ax1 = plt.subplots()  
colors = ['orange', 'lightskyblue']  
ax1.pie(values, labels=labels, autopct='%1.1f%%', shadow=True, startangle=90, color=colors)  
plt.show();
```




```
In [22]: corr = df_train.drop(columns='price_range').corr()
sns.heatmap(corr);
```



```
In [23]: target = 'price_range'
X = df_train.drop(columns=[target])
y = df_train[target]
print(f"X shape {X.shape}")
print(f"y Shape {y.shape}")
```

```
X shape (2000, 20)
y Shape (2000,)
```

```
In [25]: scalar = MinMaxScaler()
```

```
In [26]: features = X.columns
X = scalar.fit_transform(X)
```

In [27]: `print(X)`

```
[[0.22778891 0.        0.68      ... 0.        0.        1.        ]
 [0.34736139 1.        0.        ... 1.        1.        0.        ]
 [0.04141617 1.        0.        ... 1.        1.        0.        ]
 ...
 [0.94188377 0.        0.16      ... 1.        1.        0.        ]
 [0.6753507  0.        0.16      ... 1.        1.        1.        ]
 [0.00601202 1.        0.6       ... 1.        1.        1.        ]]
```

In [28]: `X_train , X_test , y_train , y_test = train_test_split(X, y , test_size=0.2 ,`

In [29]: `print("X_train shape:", X_train.shape)
print("y_train shape:", y_train.shape)
print("X_test shape :", X_test.shape)
print("y_test shape:", y_test.shape)`

```
X_train shape: (1600, 20)
y_train shape: (1600,)
X_test shape : (400, 20)
y_test shape: (400,)
```

In [30]: `model_lr = LogisticRegression(max_iter=1000)`

In [31]: `model_lr.fit(X_train,y_train)`

Out[31]: `LogisticRegression(max_iter=1000)`


In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

In [32]: `training_acc_lr= model_lr.score(X_train , y_train)
print(f"Training accuracy: {training_acc_lr}")`

```
Training accuracy: 0.938125
```

```
In [33]: ConfusionMatrixDisplay.from_estimator(model_lr, X_test, y_test);
```



```
In [34]: model_svc= SVC()
```

```
In [36]: model_svc.fit(X_train, y_train)
```

Out[36]: SVC()


In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [37]: training_acc_svc = model_svc.score(X_train , y_train)
print(f"Testing accuracy: {training_acc_svc}")
```

Testing accuracy: 0.969375

```
In [38]: ConfusionMatrixDisplay.from_estimator(model_svc, X_test, y_test);
```



```
In [51]: models = pd.DataFrame({
    "Models": ["Logistic Regression" , "SVM"],
    "Score": ["testing_acc_lr" , "testing_acc_svc"]
})
models.sort_values(by="Score" , ascending=False)
```

```
Out[51]:
```


	Models	Score
1	SVM	testing_acc_svc
0	Logestic Regression	testing_acc_lr

```
In [58]: importances = model_lr.coef_[0]


odds_ratios = pd.Series(np.exp(importances) , index = features).sort_values()
odds_ratios.head()
```

```
Out[58]: ram          7.023289e-07
battery_power  3.356091e-02
px_height     7.997938e-02
px_width      1.317819e-01
int_memory    8.198506e-01
dtype: float64
```

```
In [59]: odds_ratios.tail().plot(kind= 'barh')
plt.xlabel("Odds Ratio");
```



```
In [60]: odds_ratios.head().plot(kind= 'barh')
plt.xlabel("Odds Ratio");
```



Thanks!